

Drift-Sense Phase 2 — Failure Analysis

PS-02 · Applied Materials · Team "The T Guys" — Aryan Chourasia, Govinda Podder, Ashish Bajaj, Devaansh Gupta · Thapar Institute of Engineering & Technology · 3 September 2026

Every number below is labelled by source, because the two disagree. [REAL] = the organizers' 20-pair sample, scored against their withheld ground truth — never trained on, never used to select a threshold, checkpoint or signal. [OURS] = our generator, or their generator source re-run with new seeds. Twice, a change our own data endorsed was *rejected* when checked against [REAL]; both are documented in §4.

1. Position, and the arc that produced it

The shipped configuration is classical-led: solve.py (NCC + scale/rotation search + parabolic sub-pixel refinement) supplies all six output columns. The learned matcher remains wired into the declared DriftRoute architecture but is off by default — **because it was measured to lose, not assumed to** (§4.1).

Configuration, in order	Loc. /40 [REAL]	Change
Original Phase 2 build (net x,y; FOUND_PEAK 0.68)	13.35	—
+ classical x,y instead of net	23.10	Net overfit our textures
+ FOUND_PEAK 0.68 → 0.53	29.70	Cost-asymmetry recalibration
+ p008 rotation-aware scale scan	35.60	Root-caused bug (§2.1)

[REAL] localization **35.60/40** (Set A 1.000, Set B 0.800), 15/16 present within 5 px, rejection F1 **0.968**, rotation **10/10**, scale **8.62/10**, calibration AUC 0.789, CPU-only 1.3–3.5 s/pair against a 5 s budget. [OURS, their generator re-run, 350 pairs over two disjoint draws] rubric core **71.0–73.5/85**. We expect the blind 200 nearer the second figure: the sample's Set B is milder than the real one, and the organizers' own naive baseline scores 0.80 mean credit on it — above the 0.30–0.55 difficulty band their specification targets.

2. Failures found by measurement, root-caused, and fixed

2.1 p008 — we failed a pair the organizers' own baseline solved

The most instructive single failure. On p008 the provided ZNCC baseline earned full credit and we earned zero. Tracing the ranking rather than dismissing it as noise: the scale search evaluated every candidate zoom **at angle 0 only**, while the pipeline searched $\pm 5^\circ$. On a rotated pair the wrong zoom therefore won, and every later stage inherited it. One-line fix (SCAN_ANGLES = PHASE2_ANGLES). **Effect: 29.70 → 35.60/40 [REAL]**, with Set A rising to a perfect 1.000. The lesson generalises: a coarse stage that marginalises over the wrong variable silently caps everything downstream.

2.2 A "fundamental limit" that turned out to be our own bug

For several days we believed absent-pair rejection had a *signal separability ceiling* — absent pairs reached peak NCC 0.967 against a present median of 0.933, and no threshold separated them. We had written it up as an inherent limitation. **It was a defect in our generator:** absent search images were rendered from the reference's own layout with the landmark removed, so the two periodic backgrounds were byte-identical and a periodic template matched the decoy perfectly. Rendering absent widens from an independently instantiated, same-family layout whose lattice does not co-register:

Absent-pair signal [OURS, seed 950000]	Before	After
Peak NCC — median / maximum	0.822 / 0.967	0.533 / 0.877
Rejection AUC (raw peak) · F1	0.789 · 0.876	0.945 · 0.939

The present-pair distribution was unchanged, confirming the fix touched only the defect. **A "fundamental limit" visible only in your own data is a bug until proven otherwise.**

2.3 Two further defects, both caught by testing rather than assumption

- Rotation sign inverted.** THETA_SIGN was first chosen by pattern-matching the geometry. Correlating recovered against signed truth gave **−0.95** (median error 3.35°), and **+0.95** (median 0.26°) after flipping — since confirmed against their ground truth, their generator source and their written spec. Now **10/10 [REAL]**.
- Pose refinement overshot disclosed bounds.** Golden-section converges past an interval endpoint, so three scales fell outside [8,12] (p001 → 7.882 vs 8.00; p013 → 12.174 vs 12.00), each losing a tier — systematic, since their spec *requires* the set to reach both endpoints. Clamping the **reported** pose (not the search, which could move x,y) gave scale **8.00 → 8.62/10** with localization and rejection byte-identical.

3. Open limitations, quantified rather than hidden

3.1 The severity-4 rejection cliff — our dominant remaining loss

At the highest degradation level, dose and noise crush the true-match peak below FOUND_PEAK = 0.53, so the pair is declared absent — forfeiting *both* its localization and pose credit even though the predicted coordinate was accurate.

False-rejects among present pairs [OURS, their generator]	200-draw	150-draw
Set B severity 3	8 / 28 (29%)	4 / 23 (17%)
Set B severity 4	17 / 19 (89%)	17 / 18 (94%)
All present pairs	25 / 156 (16%)	22 / 110 (20%)

Localization at severity 4 stays sub-pixel; the credit is lost purely to the presence flag. The cause is structural: **one scalar threshold cannot simultaneously reject periodic absent decoys (wants a high threshold) and accept crushed present peaks (wants a low one)** — the

distributions overlap. The honest fix is a severity-aware presence signal, not a different scalar. We deliberately did **not** retune the threshold against the 20 validation pairs.

3.2 Generator fidelity — the limitation underneath the others

Our present-pair degradation is still too mild: ~4% of our present pairs fall below peak 0.55 against roughly 50% in the organizers' sample. Our synthetic errors are dominated by confident mislocalizations on periodic structure; theirs by low-confidence rejections. **That mismatch is why a signal tuned on our data can invert on theirs** (§4.2), and it is what we would fix first with more time. Calibration AUC — 0.789 [REAL] / 0.740 [OURS] — remains our weakest scored bucket.

4. Negative results — built, measured, and not shipped

4.1 The overfitting hypothesis, tested by removing its cause

Our most-cited assumption was "the CNN overfit our generator." We tested it directly: re-trained the identical architecture and hyperparameters on ~3,000 pairs synthesised from the *organizers' own generator source* with new seeds (their 20 pairs never trained on), selected the checkpoint on a disjoint held-out set, and scored once with weights frozen.

Configuration	Loc. /40 [REAL]	Site found ≤5 px	Set A credit
Classical (shipped)	35.60	15 / 16	1.000
CNN trained on our generator	13.35	7 / 16	0.375
CNN trained on their generator	28.78	15 / 16	0.825

The diagnosis is confirmed with a large effect — changing only the training distribution more than doubled the real-data score and took site recall from 7/16 to 15/16. **Yet classical still won, for a specific reason:** the retrained CNN matches classical on *recall* and loses on *precision* — it lands inside 5 px but outside 1 px, forfeiting tier credit that parabolic sub-pixel refinement collects. Because localization is tiered, "same recall, worse precision" costs 6.8/40. Not adopted.

4.2 A calibration signal that gained 0.21 AUC on our data and lost 0.26 on theirs

AUC is invariant to monotonic rescaling, so improving calibration needs a better *ordering*, not a recalibration curve. We tested replacing raw peak NCC with $\text{distinct} = \text{peak} \cdot (1 - \text{second_ratio})$ across three independent 150-pair draws (450 pairs, disjoint seeds), requiring a win on *all three* to guard against overfitting.

Calibration AUC	Draw 1	Draw 2	Draw 3	Organizers' 20 [REAL]
Raw peak NCC (shipped)	0.603	0.561	0.531	0.789
distinct (candidate)	0.786	0.745	0.785	0.526

Large, consistent gains on our generator; a **0.26 regression on real data**. The cause is the mismatch in §3.2 — *distinct* catches confident mislocalizations, which our generator over-produces and the real sample barely contains, while raw peak better orders the low-confidence rejection that is the real sample's only error. **Rejected; raw peak retained.** Validating only on our own data would have shipped a change that lowered our score. A Fourier–Mellin invariant pose estimator was rejected on the same principle: ~90% median scale error against the grid search's 1.1%, because the 1 μm reference field holds only 3–14 lattice periods — too few for a spectral peak.

5. A process failure, found in final verification

Our requirements.txt pinned the exact versions from pip freeze on our Python 3.12 development machine. The reference machine is **Python 3.11**, and two of those releases are 3.12-only (numpy 2.5.1 declares Requires-Python >=3.12). pip install -r requirements.txt therefore **failed outright** on a clean 3.11 environment — a defect that would have scored zero on every pair while every accuracy figure in this document remained true. It was invisible because all prior verification ran on 3.12. Replaced with ranges and re-verified end-to-end on Python 3.11.9: install succeeds and the prediction file is byte-identical to the 3.12 run. **An environment assumption is a failure mode, and must be tested on the target version.**

6. What we would do next

A severity-aware presence decision (§3.1) is the largest recoverable loss — combine peak with a degradation estimate so the accept threshold adapts, instead of one scalar serving both regimes. That depends on **closing the generator-fidelity gap** (§3.2), since no signal tuned on our data can be trusted until our error population matches theirs. Beyond those, their specification requires absent decoys to carry macro-structure the search image lacks, so a rejection signal keyed on **large-scale structure** would address the cause rather than tune a threshold.

Validation discipline. The organizers' 20 sample pairs were scored only — never trained on, never used to select a threshold, checkpoint or signal. Every adopted change was justified on our own data or from the disclosed specification, then *checked* against the sample; the two changes our data endorsed but the sample refuted (§4.1, §4.2) were rejected. Datasets regenerate from documented seeds; the shipped path needs only numpy, pillow and scipy — no GPU, no torch, no network.